

CMSC 471

Spring 2014

Class #10

Thursday, February 27, 2014
Probabilistic Reasoning

Professor Marie desJardins, mariedj@cs.umbc.edu

Today's Class

- Probability theory
- Bayesian inference
 - From the joint distribution
 - Using independence/factoring
 - From sources of evidence

Bayesian Reasoning

Chapter 13

Sources of Uncertainty

- Uncertain **inputs**
 - Missing data
 - Noisy data
 - Uncertain **knowledge**
 - Multiple causes lead to multiple effects
 - Incomplete enumeration of conditions or effects
 - Incomplete knowledge of causality in the domain
 - Probabilistic/stochastic effects
 - Uncertain **outputs**
 - Abduction and induction are inherently uncertain
 - Default reasoning, even in deductive fashion, is uncertain
 - Incomplete deductive inference may be uncertain
- Probabilistic reasoning only gives probabilistic results
(summarizes uncertainty from various sources)


Decision Making with Uncertainty

- **Rational** behavior:

- For each possible action, identify the possible outcomes
- Compute the **probability** of each outcome
- Compute the **utility** of each outcome
- Compute the probability-weighted **(expected) utility** over possible outcomes for each action
- Select the action with the highest expected utility (principle of **Maximum Expected Utility**)

Why Probabilities Anyway?

- Kolmogorov showed that three simple axioms lead to the rules of probability theory
 - De Finetti, Cox, and Carnap have also provided compelling arguments for these axioms
- 1. All probabilities are between 0 and 1:
 - $0 \leq P(a) \leq 1$
- 2. Valid propositions (tautologies) have probability 1, and unsatisfiable propositions have probability 0:
 - $P(\text{true}) = 1$; $P(\text{false}) = 0$
- 3. The probability of a disjunction is given by:
 - $P(a \vee b) = P(a) + P(b) - P(a \wedge b)$



Probability Theory

- **Random variables**
 - Domain
- **Atomic event**: complete specification of state
- **Prior probability**: degree of belief without any other evidence
- **Joint probability**: matrix of combined probabilities of a set of variables
- Alarm, Burglary, Earthquake
 - Boolean (like these), discrete, continuous
- $\text{Alarm}=\text{True} \wedge \text{Burglary}=\text{True} \wedge \text{Earthquake}=\text{False}$
 $\text{alarm} \wedge \text{burglary} \wedge \neg \text{earthquake}$
- $P(\text{Burglary}) = .1$
- $P(\text{Alarm}, \text{Burglary}) =$

	alarm	\neg alarm
burglary	.09	.01
\neg burglary	.1	.8

Probability Theory: Definitions

- **Conditional probability:**
probability of effect given causes
- **Computing conditional prob:**
 - $P(a | b) = P(a \wedge b) / P(b)$
 - $P(b)$: **normalizing** constant
- **Product rule:**
 - $P(a \wedge b) = P(a | b) P(b)$
- **Marginalizing:**
 - $P(B) = \sum_a P(B, a)$
 - $P(B) = \sum_a P(B | a) P(a)$
(**conditioning**)

Try It...

	alarm	\neg alarm
burglary	.09	.01
\neg burglary	.1	.8

- $P(\text{alarm} \mid \text{burglary}) = ??$
- $P(\text{burglary} \mid \text{alarm}) = ??$
- $P(\text{burglary} \wedge \text{alarm}) = ??$
- $P(\text{alarm}) = ??$

- **Computing conditional prob:**

- $P(a \mid b) = P(a \wedge b) / P(b)$
- $P(b)$: **normalizing** constant

- **Product rule:**

- $P(a \wedge b) = P(a \mid b) P(b)$

- **Marginalizing:**

- $P(B) = \sum_a P(B, a)$
- $P(B) = \sum_a P(B \mid a) P(a)$
(**conditioning**)

Probability Theory (cont.)

- **Conditional probability:**
probability of effect given causes
- **Computing conditional probs:**
 - $P(a | b) = P(a \wedge b) / P(b)$
 - $P(b)$: **normalizing** constant
- **Product rule:**
 - $P(a \wedge b) = P(a | b) P(b)$
- **Marginalizing:**
 - $P(B) = \sum_a P(B, a)$
 - $P(B) = \sum_a P(B | a) P(a)$
(**conditioning**)
- $P(\text{burglary} | \text{alarm}) = .47$
 $P(\text{alarm} | \text{burglary}) = .9$
- $P(\text{burglary} | \text{alarm}) =$
 $P(\text{burglary} \wedge \text{alarm}) / P(\text{alarm})$
 $= .09 / .19 = .47$
- $P(\text{burglary} \wedge \text{alarm}) =$
 $P(\text{burglary} | \text{alarm}) P(\text{alarm}) =$
 $.47 * .19 = .09$
- $P(\text{alarm}) =$
 $P(\text{alarm} \wedge \text{burglary}) +$
 $P(\text{alarm} \wedge \neg \text{burglary}) =$
 $.09 + .1 = .19$

Example: Inference from the Joint

	alarm		\neg alarm	
	earthquake	\neg earthquake	earthquake	\neg earthquake
burglary	.01	.08	.001	.009
\neg burglary	.01	.09	.01	.79

$$\begin{aligned}
 P(\text{Burglary} \mid \text{alarm}) &= \alpha P(\text{Burglary}, \text{alarm}) \\
 &= \alpha [P(\text{Burglary}, \text{alarm}, \text{earthquake}) + P(\text{Burglary}, \text{alarm}, \neg \text{earthquake})] \\
 &= \alpha [(.01, .01) + (.08, .09)] \\
 &= \alpha [(.09, .1)]
 \end{aligned}$$

Since $P(\text{burglary} \mid \text{alarm}) + P(\neg \text{burglary} \mid \text{alarm}) = 1$, $\alpha = 1/(.09+.1) = 5.26$
 (i.e., $P(\text{alarm}) = 1/\alpha = .19$ – **quizlet**: how can you verify this?)

$$P(\text{burglary} \mid \text{alarm}) = .09 * 5.26 = .474$$

$$P(\neg \text{burglary} \mid \text{alarm}) = .1 * 5.26 = .526$$

Exercise: Inference from the Joint

$p(\text{smart} \wedge \text{study} \wedge \text{prep})$	smart		$\neg\text{smart}$	
	study	$\neg\text{study}$	study	$\neg\text{study}$
prepared	.432	.16	.084	.008
$\neg\text{prepared}$.048	.16	.036	.072

- Queries:
 - What is the prior probability of *smart*?
 - What is the prior probability of *study*?
 - What is the conditional probability of *prepared*, given *study* and *smart*?
- Save these answers for later! ☺

Independence

- When two sets of propositions do not affect each others' probabilities, we call them **independent**, and can easily compute their joint and conditional probability:
 - Independent $(A, B) \Leftrightarrow P(A \wedge B) = P(A) P(B)$, $P(A | B) = P(A)$
- For example, {moon-phase, light-level} might be independent of {burglary, alarm, earthquake}
 - Then again, it might not: Burglars might be more likely to burglarize houses when there's a new moon (and hence little light)
 - But if we know the light level, the moon phase doesn't affect whether we are burglarized
 - Once we're burglarized, light level doesn't affect whether the alarm goes off
- We need a more complex notion of independence, and methods for reasoning about these kinds of relationships

Exercise: Independence

$p(\text{smart} \wedge \text{study} \wedge \text{prep})$	smart		$\neg\text{smart}$	
	study	$\neg\text{study}$	study	$\neg\text{study}$
prepared	.432	.16	.084	.008
$\neg\text{prepared}$.048	.16	.036	.072

- Queries:
 - Is *smart* independent of *study*?
 - Is *prepared* independent of *study*?

Conditional Independence

- Absolute independence:
 - A and B are **independent** if $P(A \wedge B) = P(A) P(B)$; equivalently, $P(A) = P(A | B)$ and $P(B) = P(B | A)$
- A and B are **conditionally independent** given C if
 - $P(A \wedge B | C) = P(A | C) P(B | C)$
- This lets us decompose the joint distribution:
 - $P(A \wedge B \wedge C) = P(A | C) P(B | C) P(C)$
- Moon-Phase and Burglary are ***conditionally independent given*** Light-Level
- Conditional independence is weaker than absolute independence, but still useful in decomposing the full joint probability distribution

Exercise: Conditional Independence

$p(\text{smart} \wedge \text{study} \wedge \text{prep})$	smart		$\neg\text{smart}$	
	study	$\neg\text{study}$	study	$\neg\text{study}$
prepared	.432	.16	.084	.008
$\neg\text{prepared}$.048	.16	.036	.072


- Queries:
 - Is *smart* conditionally independent of *prepared*, given *study*?
 - Is *study* conditionally independent of *prepared*, given *smart*?

Bayes's Rule

- Bayes's rule is derived from the product rule:
 - $P(Y | X) = P(X | Y) P(Y) / P(X)$
- Often useful for diagnosis:
 - If X are (observed) effects and Y are (hidden) causes,
 - We may have a model for how causes lead to effects ($P(X | Y)$)
 - We may also have prior beliefs (based on experience) about the frequency of occurrence of effects ($P(Y)$)
 - Which allows us to reason abductively from effects to causes ($P(Y | X)$).

Bayesian Inference

- In the setting of diagnostic/evidential reasoning



- Know prior probability of hypothesis $P(H_i)$
- conditional probability $P(E_j | H_i)$
- Want to compute the *posterior probability* $P(H_i | E_j)$
- Bayes's theorem (formula 1):

$$P(H_i | E_j) \propto P(H_i)P(E_j | H_i) / P(E_j)$$

Simple Bayesian Diagnostic Reasoning

- Knowledge base:
 - Evidence / manifestations: E_1, \dots, E_m
 - Hypotheses / disorders: H_1, \dots, H_n
 - E_j and H_i are **binary**; hypotheses are **mutually exclusive** (non-overlapping) and **exhaustive** (cover all possible cases)
 - Conditional probabilities: $P(E_j | H_i)$, $i = 1, \dots, n$; $j = 1, \dots, m$
- Cases (evidence for a particular instance): E_1, \dots, E_l
- Goal: Find the hypothesis H_i with the highest posterior
 - $\text{Max}_i P(H_i | E_1, \dots, E_l)$

Bayesian Diagnostic Reasoning II

- Bayes' rule says that
 - $P(H_i | E_1, \dots, E_l) = P(E_1, \dots, E_l | H_i) P(H_i) / P(E_1, \dots, E_l)$
- Assume each piece of evidence E_i is conditionally independent of the others, *given* a hypothesis H_i , then:
 - $P(E_1, \dots, E_l | H_i) = \prod_{j=1}^l P(E_j | H_i)$
- If we only care about relative probabilities for the H_i , then we have:
 - $P(H_i | E_1, \dots, E_l) = \alpha P(H_i) \prod_{j=1}^l P(E_j | H_i)$

Limitations of Simple Bayesian Inference

- Cannot easily handle multi-fault situations, nor cases where intermediate (hidden) causes exist:
 - Disease D causes syndrome S, which causes correlated manifestations M_1 and M_2
- Consider a composite hypothesis $H_1 \wedge H_2$, where H_1 and H_2 are independent. What is the relative posterior?
 - $$\begin{aligned} P(H_1 \wedge H_2 \mid E_1, \dots, E_l) &= \alpha P(E_1, \dots, E_l \mid H_1 \wedge H_2) P(H_1 \wedge H_2) \\ &= \alpha P(E_1, \dots, E_l \mid H_1) P(H_1) P(H_2) \\ &= \alpha \prod_{j=1}^l P(E_j \mid H_1 \wedge H_2) P(H_1) P(H_2) \end{aligned}$$
- How do we compute $P(E_j \mid H_1 \wedge H_2)$??

Limitations of Simple Bayesian Inference II

- Assume H_1 and H_2 are independent, given E_1, \dots, E_l ?
 - $P(H_1 \wedge H_2 | E_1, \dots, E_l) = P(H_1 | E_1, \dots, E_l) P(H_2 | E_1, \dots, E_l)$
- This is a very unreasonable assumption
 - Earthquake and Burglar are independent, but *not* given Alarm:
 - $P(\text{burglar} | \text{alarm}, \text{earthquake}) \ll P(\text{burglar} | \text{alarm})$
- Another limitation is that simple application of Bayes's rule doesn't allow us to handle causal chaining:
 - A: this year's weather; B: cotton production; C: next year's cotton price
 - A influences C indirectly: $A \rightarrow B \rightarrow C$
 - $P(C | B, A) = P(C | B)$
- Need a richer representation to model interacting hypotheses, conditional independence, and causal chaining
- Next time: conditional independence and Bayesian networks!